



Building Resilient Data Pipelines

An introduction to Building Data Pipelines in Go on Kubernetes

Grant Griffiths
Software Engineer
Platform Cloud Engineering
GE Digital

May 16th 2018

PREDIX

Who am I?

Senior Software Engineer

GE Digital - Platform Cloud Engineering

Organizer for the GE Go User Group



@ggriffiths




@griffithsgrant

Climbing + Mountaineering



What we'll cover


1. Industrial IoT and Data @ GE
2. Introduction to Data Pipelines
3. Sample Data Pipeline in Go
4. Resiliency Testing our Data Pipeline
5. Results and takeaways



1. Industrial IoT and Data @ GE
2. Introduction to Data Pipelines
3. Sample Data Pipeline in Go
4. Reliability Testing our Data Pipeline
5. Results and takeaways



Who here has owned a GE Fridge, Lightbulb, or Microwave?



Actually, we make
HUGE Industrial assets...



MAX CRANE STRUCTURE CAP 12.5t





Industrial Internet of Things

- GE Assets produce a ton of useful data
- Valuable for gaining insights into these assets
- Asset Performance Management (APM)
- Small percent increase in efficiency saves billions of dollars



Predix

- Multi-cloud platform for the Industrial Internet of Things
- Many services and applications optimized for industrial data

PREDIX



Some Customers

- Schindler
- Exelon
- Rosneft
- BP
- GE Power
- GE Aviation
- GE Renewables




Schindler



Exelon[®]



Monitoring and Diagnostic Architecture




Data Pipeline component (Java version)

- **Stack:**
 - Stateful Java app
 - Java and Apache Apex
- **Largest deployment (out of several)**
 - 150 Cassandra nodes
 - 30 Kafka nodes
 - 144 Apache Apex containers
- **Results:**
 - 900,000 C* writes/second (peak)




Data Pipeline component (Go version)

- **Stack:**
 - Stateless Go app
 - Kubernetes & Cloud Foundry (prod environment specific)
- **Prod Configuration:**
 - Cloud Foundry application
 - 8+ instances
- **Results:**
 - Similar performance at lower scale (smaller dedicated cluster)
 - Pushing to multiple other dedicated environments soon!



Rewrite motivations



- Change in service vision/purpose
 - Originally – wanted customer specific data models/parsing
 - Now – standard data model, parsing
- Operational cost (\$\$\$)
 - Managing a Hadoop cluster
 - Resources (RAM/CPU/Disk)
- Moving towards Kubernetes
 - Simple Go microservice
- We love Go!

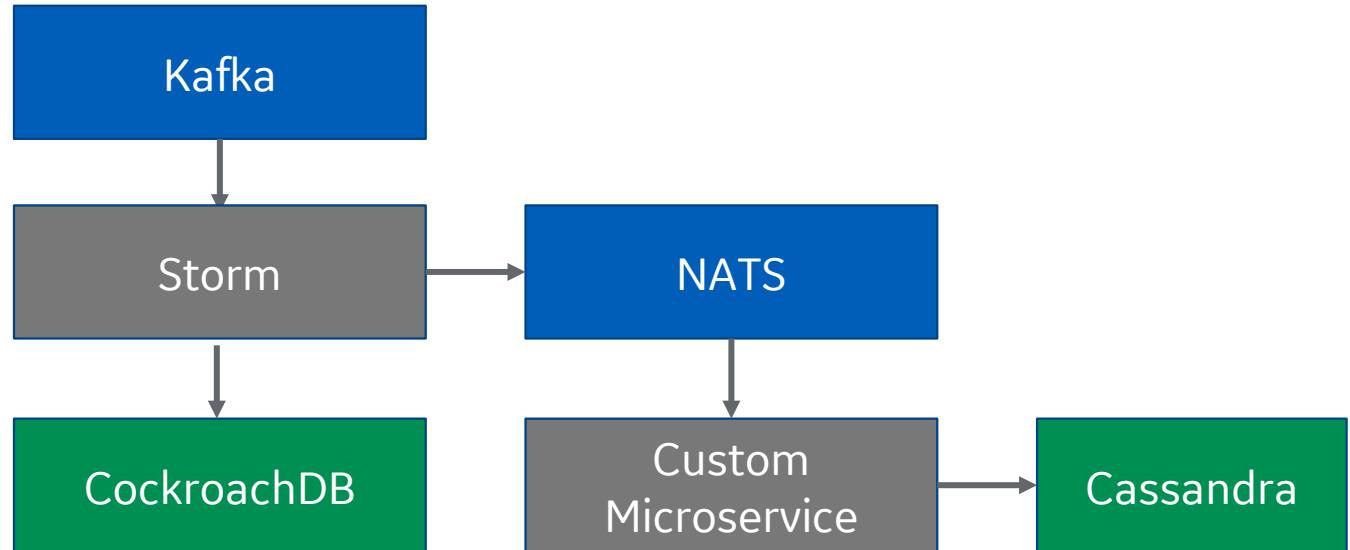
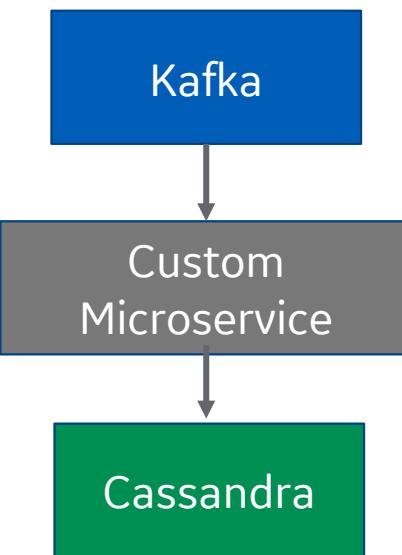


1. Data @ GE
2. **Introduction to Data Pipelines**
3. Sample Data Pipeline in Go
4. Reliability Testing our Data Pipeline
5. Results and takeaways



Introduction to Data Pipelines

- Move data from one system to another
- Performing transformations and business logic

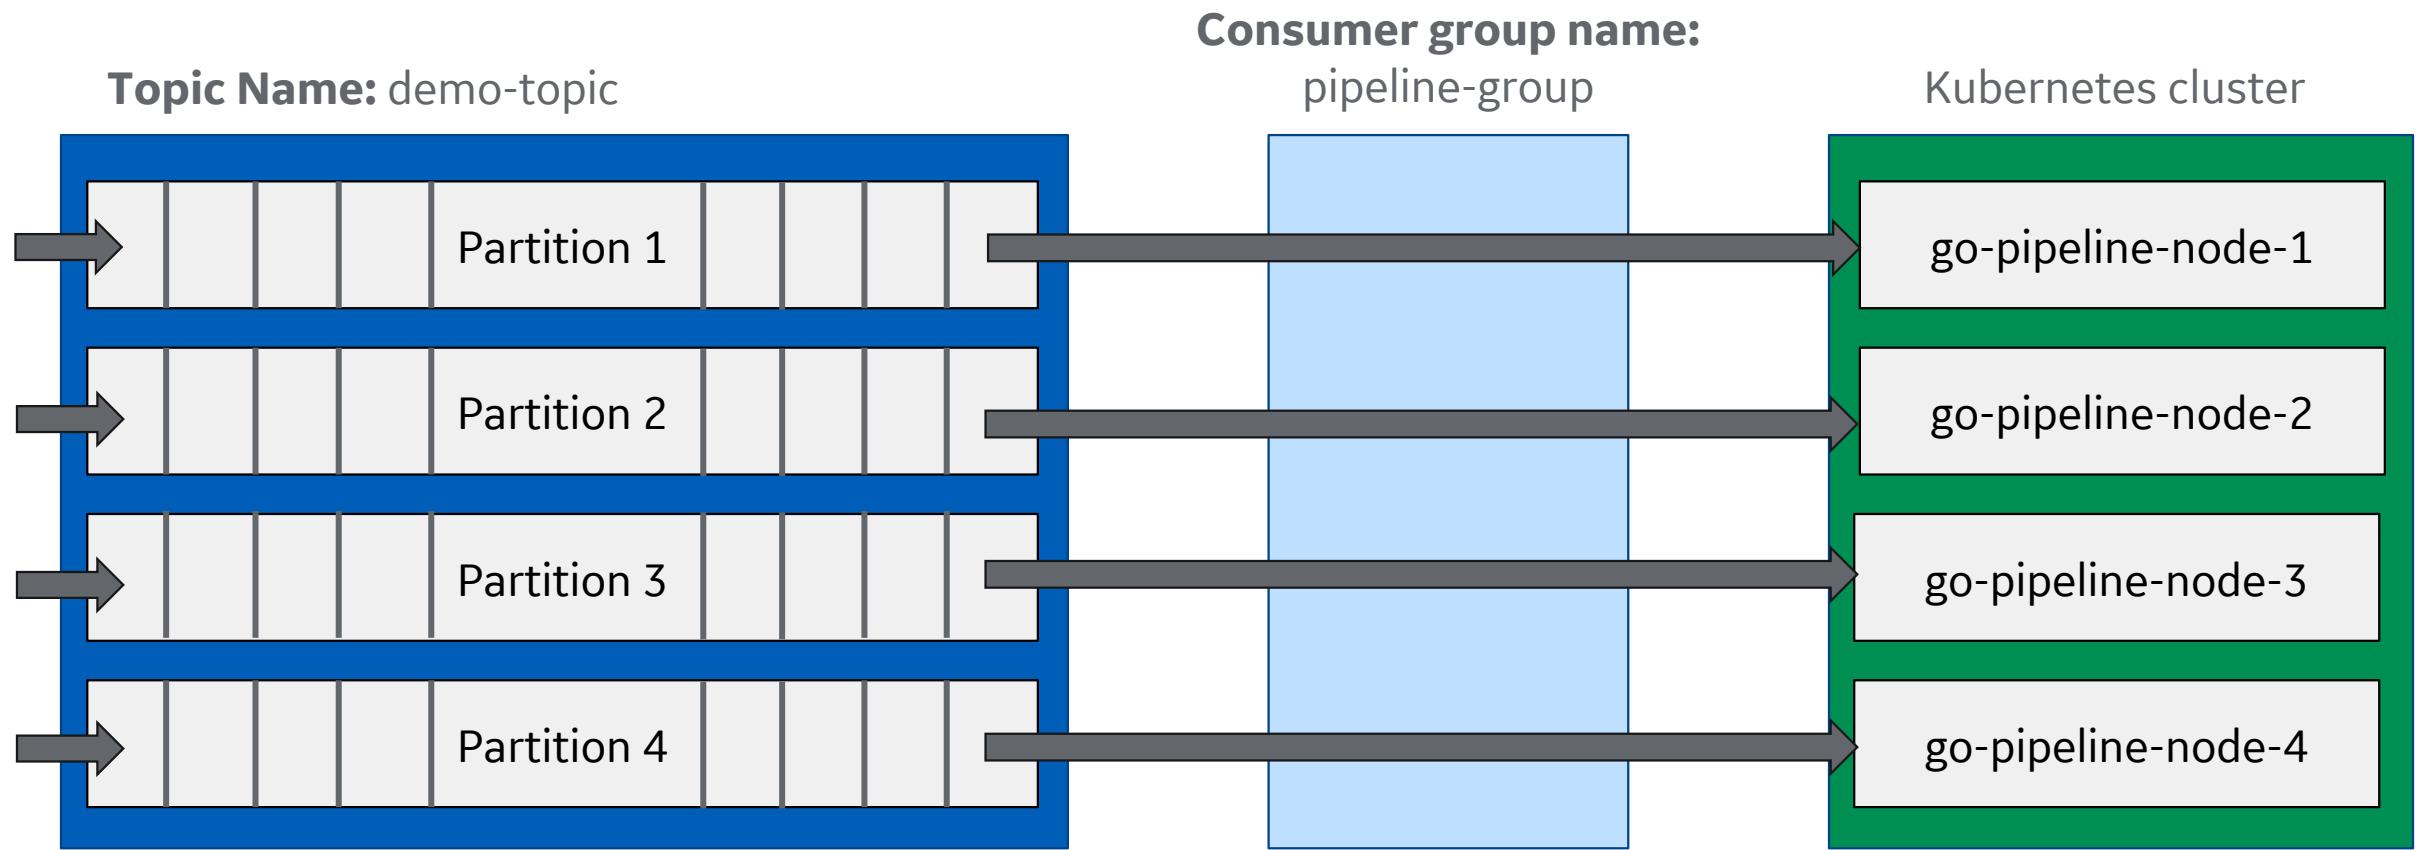


Data source: Apache Kafka


- Publish/Subscribe messaging system
- Parallelized with topic partitions
- High throughput
- Very widely used
- Java, open source - github.com/apache/kafka



Consumer groups Example 1: 1:1



Consumer groups Example: 2:1



Using Kafka with Go

- Many libraries
 - github.com/Shopify/sarama + github.com/bsm/sarama-cluster
 - github.com/confluentinc/confluent-kafka-go
 - github.com/segmentio/kafka-go (June 2017)
- Chose Sarama + Sarama Cluster
 1. No CGo dependency
 2. Most mature library (at the time)
 3. Wrote internal tooling + documentation around it for ease of use
- Pick what works for you



Data store: Apache Cassandra

- Column-oriented database
- Fault Tolerant - replicated
- Scalable
 - Apple: over 75,000 nodes storing over 10 PB of data
 - Netflix: 2,500 nodes, 420 TB, 1 trillion requests
 - GE: 1,250 nodes, find size, find throughput
- Java, open source github.com/apache/cassandra



Go and Cassandra


- github.com/gocql/gocql
- For high performance data bindings:
 - github.com/scylladb/gocqlx



1. Data @ GE
2. Introduction to Data Pipelines
- 3. Sample Data Pipeline in Go**
4. Reliability Testing our Data Pipeline
5. Results and takeaways



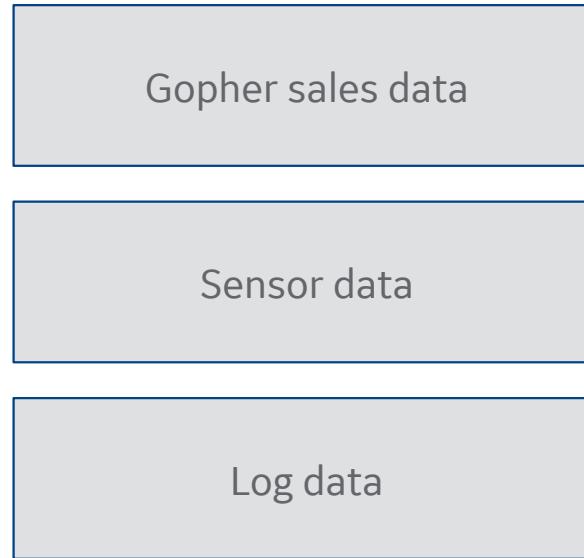
By show of hands...
Who likes the Go Gopher?



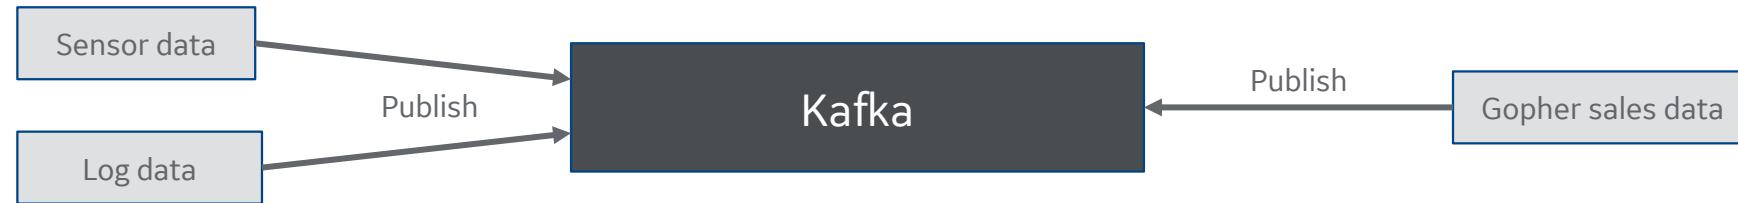
I bring to you... Gophers “R” Us



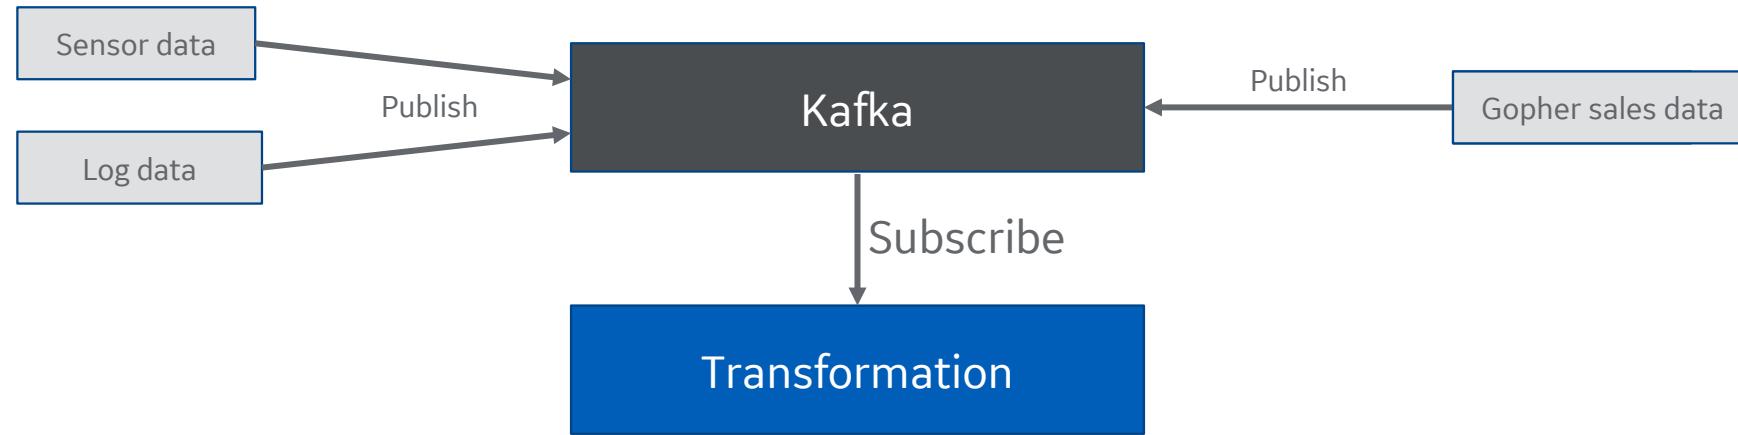
Example: Data Pipeline (Gophers “R” Us)




Data pipeline architecture



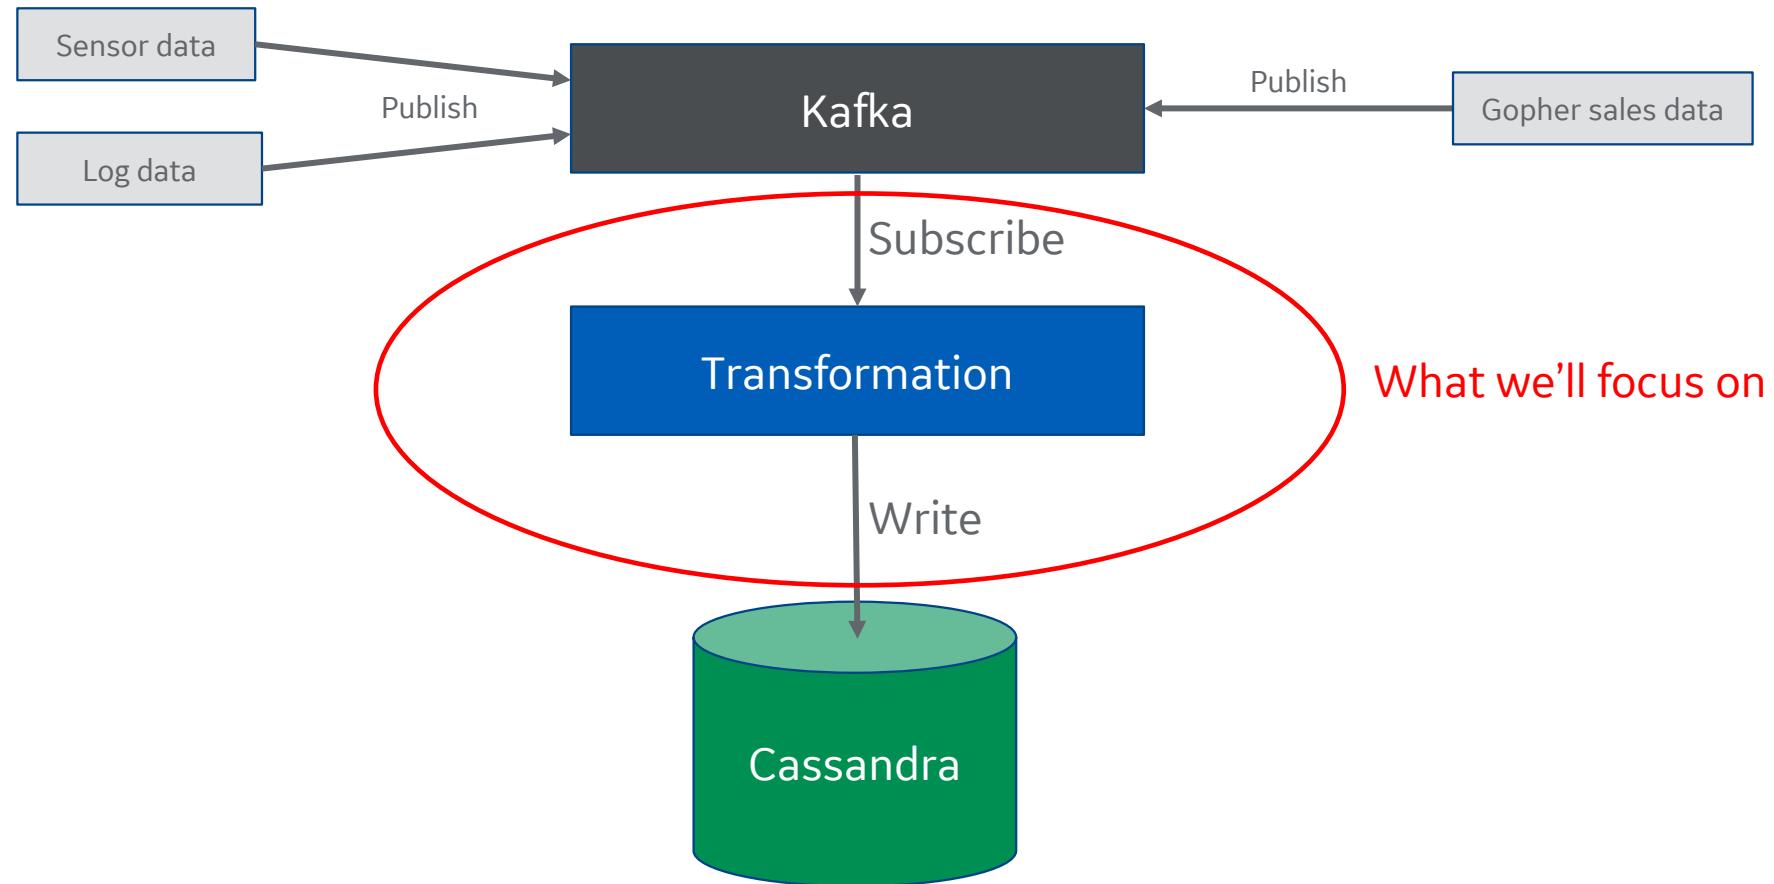
Data pipeline architecture



Data pipeline architecture




Data pipeline architecture




Simplified Application Flow: 3 Easy Steps

```
for {  
    select {  
        case msg := <-consumer.Messages():  
            event := Transform(&msg)  
            sink.Write(&event)  
    }  
}
```



Subscribing to Kafka

```
for {  
    select {  
        case msg := <-consumer.Messages():  
            event := Transform(&msg)  
            sink.Write(&event)  
    }  
}
```



Subscribing to Kafka

```
for {
    select {
        case msg, ok := <-consumer.Messages():
            log.Printf("Messages received: %v", *msg)

        case notification, ok := <-consumer.Notifications():
            log.Printf("Rebalance received: %v", *notification)

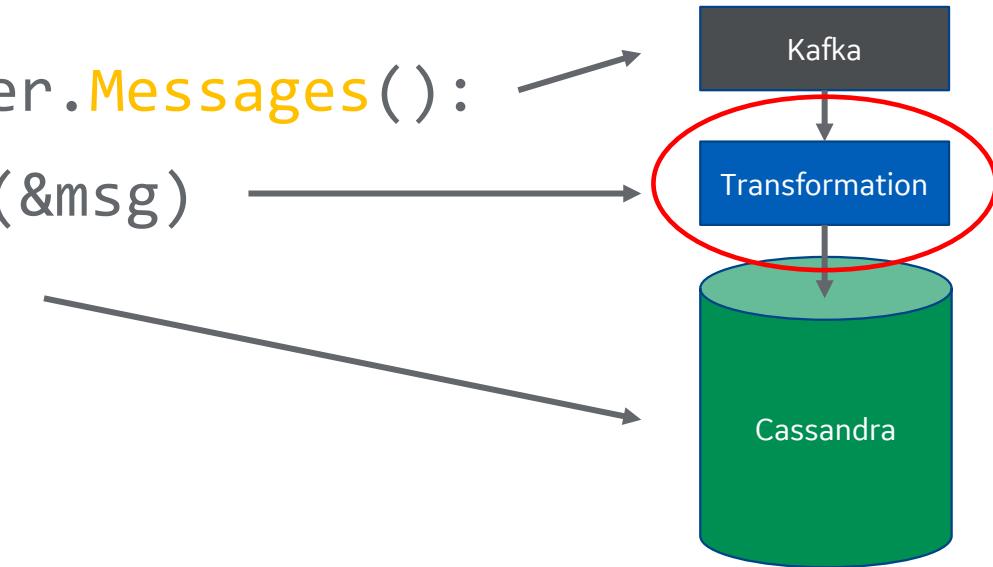
        case err, ok := <-consumer.Errors():
            log.Printf("Error received: %s", err.Error())
    }
}
```

1. **Messages channel:** Data from Kafka
2. **Notifications channel:** Rebalance notifications
3. **Errors channel:** Errors in offset management



Handling messages

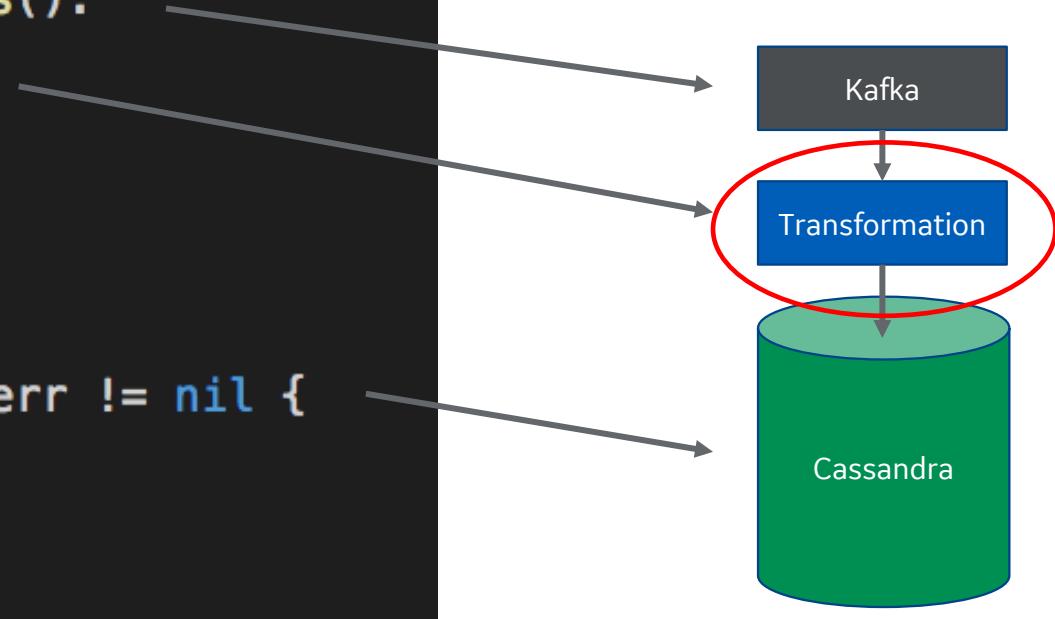
```
for {  
    select {  
        case msg := <-consumer.Messages():  
            event := Transform(&msg)  
            sink.Write(&event)  
    }  
}
```



Handling messages

```
case msg, _ := <-consumer.Messages():
    event, err := Transform(msg)
    if err != nil {
        ...
    }

    if err := sink.Write(event); err != nil {
        ...
    }
    consumer.MarkOffset(msg, "")
```



Transforming our message

Gophers “R” Us complicated business logic

```
// Transform take a sarama Message and transforms it based
// on our business logic
func Transform(msg *sarama.ConsumerMessage) (*StoreEvent, error) {
    var event StoreEvent
    if err := json.Unmarshal(msg.Value, &event); err != nil {
        ...
    }
    event.AssignDepartment()

    return &event, nil
}
```



Transforming our message

Gophers “R” Us complicated business logic

```
// AssignDepartment assigns a department for an item based on it's name
func (se *StoreEvent) AssignDepartment() {
    switch se.ItemName {
        case "blue-gopher", "pink-gopher", "purple-gopher":
            se.ItemDepartment = deptGophers

        case "gophercon-us-sticker", "gosf-sticker", "gophercon-uk-sticker":
            se.ItemDepartment = deptStickers


        case "gophercon-us-shirt", "gophercon-uk-shirt":
            se.ItemDepartment = deptClothing

        default:
            se.ItemDepartment = deptMisc
    }
}
```



Writing to Cassandra

```
for {  
    select {  
        case msg := <-consumer.Messages():  
            event := Transform(&msg)  
            sink.Write(&event)  
    }  
}
```



Writing to Cassandra

```
// Write persists an event to Cassandra
func (cs *CassandraSink) Write(e *StoreEvent) error {
    stmt, names := qb.Insert(fmt.Sprintf("%s.%s", cs.Keyspace, cs.StoreEventTable)).
        Columns("name", "price", "dept").
        ToCql()
    q := gocqlx.Query(cs.Session.Query(stmt), names).Bind(e.ItemName, e.Price, e.ItemDepartment)
    return q.Exec()
}
```



Graceful shutdown

```
signals := make(chan os.Signal, 1)
for {
    select {
    case sig, _ := <-signals:
        switch sig {
        case syscall.SIGTERM:
            log.Printf("Gracefully shutting down. Goodbye!")
            consumer.Close()
            sink.Session.Close()
            os.Exit(0)
        }
    }
}
```

1. Setup channel for listening to OS Signals
2. Listen on os signals channel
3. Handle SIGTERM for graceful shutdown
 - k8s sends this signal when containers are stopped, scaled down, etc
4. Grace period (10-30s)
 - Configurable in k8s:
terminationGracePeriodSeconds
 - Container killed after this period
(or you can os.Exit(0) manually)



Deploying to Kubernetes

```
---  
kind: Deployment  
apiVersion: extensions/v1beta1  
metadata:  
  name: gophers-r-us-pipeline  
spec:  
  template:  
    metadata:  
      labels:  
        app: gophers-r-us-pipeline  
        id: "1"  
    spec:  
      containers:  
      - name: gophers-r-us-pipeline  
        image: ggriffiths/gophers-r-us-pipeline:v1  
        ports:  
        - containerPort: 8080  
      terminationGracePeriodSeconds: 15
```



1. Data @ GE
2. Introduction to Data Pipelines
3. Sample Data Pipeline in Go
- 4. Reliability Testing our Data Pipeline**
5. Results and takeaways



Reliability Testing

- Systems fail
 - How does our pipeline behave during times like these?
 - How can we remedy these failures?
 - How can we ensure customer data is not lost?
- Embrace failure scenarios
 - or they will embrace you at 3 AM
- If this interests you: Google SRE Book
landing.google.com/sre/book/index.html



Reliability Testing our Data Pipeline

- What can fail?
 - Data pipeline node(s)
 - Kafka node(s)
 - Cassandra node(s)
 - Kubernetes can fail
- How does our pipeline behave when this happens?
 - We can write a test
- How can we remedy these failures?



Reliability Test example

```
Describe("Given our Data pipeline is running successfully, reading from Kafka and writing to C*", func() {
    Context("When our Cassandra cluster goes down", func() {
        It("Then no data should be lost", func() {
            initialCount := cassandraRowCount("keyspace", "item_table")

            desiredMessageCount := 50
            for i := 0; i < desiredMessageCount; i++ {
                pushMessages("demo-topic", 1)

                // Take Cassandra down for 20 seconds
                switch i {
                    case 10:
                        pauseDockerImage("cassandra")

                    case 30:
                        unPauseDockerImage("cassandra")
                }


                time.Sleep(1 * time.Second)
            }

            finalCount := cassandraRowCount("keyspace", "item_table")
            Expect(finalCount - initialCount).To(Equal(desiredMessageCount))
        })
    })
})
```



Reliability Test example

- What else can we test?
 - Partial cluster failures
 - Full cluster failures
 - Pipeline failures
 - High load
 - etc
- Integration Test with Docker!




1. Data @ GE
2. Introduction to Data Pipelines
3. Sample Data Pipeline in Go
4. Reliability Testing our Data Pipeline
5. **Results and takeaways**



Results & Takeaways

- **Building Data Pipelines in Go**
 - Simple to get a small app up and running
 - Good enough community support for Kafka, Cassandra, etc
 - Use chan of os.Signals for graceful shutdown
- **Reliability testing**
 - Use docker to integration/reliability test
 - Understand how your system behaves during failures scenarios
- **Go Pipelines at GE Digital**
 - Replaced our existing Java pipeline for lower operational cost, simplicity, and performance



We're hiring!



PREDIX

Questions?

Or funding for Gophers "R" Us?



PREDIX